Jifan Chen

jifanchen.ai@gmail.com | jifan-chen.github.io

Industry Experience

AWS AI Lab — Santa Clara, US
Applied Scientist — *2023.06 – Present*

Kiro Science

- Lead the interactive (user–agent simulation) and prefix-based environment creation work for training the Kiro first-party code agent.
- Lead the Kiro live data analysis effort and code-agent training via on-policy distillation with live trajectories.

Amazon Q Science

- Co-lead the evaluation of the RAG system, focusing on citation evaluation of the final response (*CiteEval*).
- Lead the development of RAG systems that retrieve and answer questions over table content.

Applied Scientist Intern — *2022.05 – 2022.09*

- Mentors: Yuhao Zhang, Lan Liu.
- Improving cross-task generalization of unified table-to-text models with compositional task configurations.

Google Language — New York, US
Research Intern & Student Researcher — *2020.05 – 2020.11*

- Mentors: Daniel Andor, Livio Baldini Soares
- Robust question answering by incorporating linguistic explanations.

Education

University of Texas at Austin — Austin, Texas, US
Ph.D. in Computer Science — *2017 – 2023*

- Advisors: Greg Durrett, Eunsol Choi
- Dissertation: *Building Robust and Modular Question Answering Systems*

Fudan University — Shanghai, China
MSc in Computer Science — *2014 – 2017*

- Advisor: XuanJing Huang
- Research: Implicit Discourse Relation Detection on PDTB; Learning Word Embeddings using both Intrinsic and Extrinsic Information

Fudan University — Shanghai, China
BSc in Computer Science — *2010 – 2014*

Publications

1. Shoumik Saha, **Jifan Chen**, Sam Mayers, Sanjay Krishna Gouda, Zijian Wang, Varun Kumar. Breaking the Code: Security Assessment of AI Code Agents through Systematic Jailbreaking Attacks. *arXiv:2510.01359*, 2025.
2. Yumo Xu, Peng Qi, **Jifan Chen**, Kunlun Liu, Rujun Han, Lan Liu, Bonan Min, Vittorio Castelli, Arshit Gupta, Zhiguo Wang. CiteEval: Principle-Driven Citation Evaluation for Source Attribution. *ACL*, 2025.
3. Marc E. Canby, Xinchi Chen, Xing Niu, **Jifan Chen**, Bonan Min, Sergul Aydore, Vittorio Castelli. Benchmarking Query-Conditioned Natural Language Inference. *Findings of ACL*, 2025.
4. Zhengxuan Wu, Yuhao Zhang, Peng Qi, Yumo Xu, Rujun Han, Yian Zhang, **Jifan Chen**, Bonan Min, Zhiheng Huang. Dancing in Chains: Reconciling Instruction Following and Faithfulness in Language Models. *EMNLP*, 2024.
5. Manya Wadhwa, **Jifan Chen**, Jessy Li, Greg Durrett. Using Natural Language Explanations to Rescale Human Judgments. *COLM*, 2024.
6. **Jifan Chen**, Grace Kim, Aniruddh Sriram, Greg Durrett, Eunsol Choi. Complex Claim Verification with Evidence Retrieved in the Wild. *NAACL*, 2024.
7. **Jifan Chen**. Building Robust and Modular Question Answering Systems. *Ph.D. Dissertation*, 2023.
8. **Jifan Chen**, Yuhao Zhang, Lan Liu, Rui Dong, Xinchi Chen, Patrick Ng, William Yang Wang, Zhiheng Huang. Improving Cross-task Generalization of Unified Table-to-text Models with Compositional Task Configurations. *Findings of ACL*, 2023.
9. Venelin Kovatchev, Trina Chatterjee, Venkata S. Govindarajan, **Jifan Chen**, Eunsol Choi, Gabriella Chronis, Anubrata Das, Katrin Erk, Matthew Lease, Junyi Jessy Li, Yating Wu, Kyle Mahowald. Longhorns at DADC 2022: How Many Linguists Does It Take to Fool a Question Answering Model? A Systematic Approach to Adversarial Attacks. *DADC Workshop at EMNLP*, 2022.
10. **Jifan Chen**, Aniruddh Sriram, Eunsol Choi, Greg Durrett. Generating Literal and Implied Subquestions to Fact-check Complex Claims. *EMNLP*, 2022.
11. **Jifan Chen**, Eunsol Choi, Greg Durrett. Can NLI Models Verify QA Systems' Predictions? *Findings of EMNLP*, 2021.
12. Greg Durrett, **Jifan Chen**, Shrey Desai, Tanya Goyal, Lucas Kabela, Yasumasa Onoe, Jiacheng Xu. Contemporary NLP Modeling in Six Comprehensive Programming Assignments. *Proceedings of the Fifth Workshop on Teaching NLP*, 2021.
13. **Jifan Chen**, Greg Durrett. Robust Question Answering via Sub-part Alignment. *NAACL*, 2021.
14. **Jifan Chen**, Shih-ting Lin, Greg Durrett. Multi-hop Question Answering via Reasoning Chains. *arXiv*, 2019.
15. **Jifan Chen**, Greg Durrett. Understanding Dataset Design Choices for Multi-hop Reasoning. *NAACL*, 2019.
16. **Jifan Chen**, Greg Durrett. How to Learn (and How Not to Learn) Multi-hop Reasoning with Memory Networks. *OpenReview*, 2019.
17. **Jifan Chen**, Qi Zhang, Pengfei Liu, Xipeng Qiu, Xuanjing Huang. Implicit Discourse Relation Detection via a Deep Architecture with Gated Relevance Network. *ACL*, 2016. (Oral)
18. **Jifan Chen**, Qi Zhang, Pengfei Liu, Xuanjing Huang. Discourse Relations Detection via a Mixed Generative-Discriminative Framework. *AAAI*, 2016.
19. **Jifan Chen**, Kan Chen, Xipeng Qiu, Qi Zhang, Xuanjing Huang, Zheng Zhang. Learning Word Embeddings from Intrinsic and Extrinsic Views. *arXiv:1608.05852*, 2016.
20. **Jifan Chen**, Qi Zhang, Xuanjing Huang. Incorporate Group Information to Enhance Network Embed-

ding. *CIKM*, 2016. (Short paper)

21. Pengfei Liu, Xipeng Qiu, **Jifan Chen**, Xuanjing Huang. Deep Fusion LSTMs for Text Semantic Matching. *ACL*, 2016.
22. Pengfei Liu, Xipeng Qiu, Yaqian Zhou, **Jifan Chen**, Xuanjing Huang. Modelling Interaction of Sentence Pair with Coupled-LSTMs. *EMNLP*, 2016.
23. Yikang Shen, **Jifan Chen**, Xuanjing Huang. Bidirectional Long Short-Term Memory with Gated Relevance Network for Paraphrase Identification. *ICCPOL*, 2016.

Academic Service

- **Reviewer:** ACL Rolling Review (21-09 – 23-02); NeurIPS 2021–2023; AKBC 2021–2022; ICLR 2020–2023; ACL-IJCNLP 2021–2023; EMNLP 2021–2022; NLPCC 2021.
- **Workshop reviewer:** SRW @ ACL-IJCNLP 2021–2022; SRW @ NAACL 2019, 2021; SRW @ AACL-IJCNLP 2020; DeepLo @ EMNLP 2019.
- **Program committee:** Workshop on Machine Reading for Question Answering @ EMNLP 2021.
- **Volunteer:** ACL 2015 (Beijing); WSDM 2015 (Shanghai).
- Masters Admissions Committee member, UT Austin Computer Science Department, 2020.

Teaching Experience

- TA, C++ Programming Language, Fudan University — Spring 2016.
- TA, Natural Language Processing (CS388), UT Austin — Fall 2018, Fall 2021.